



Intro to Deep Learning with Julia

Bucharest Julia Meetup
9th of October 2019

Traffic Sign Detector

- Task:
 - Create a traffic sign detection library which can be used standalone or interfaced with smartphone Android app
- Constraints:
 - Localize and classify traffic signs in a series of video frames
 - 1fps processing speed
 - Run on server CPU

Traffic Sign Detector Cookbook

- Ingredients:
 - Deep learning library:
 - For defining and training a model and/or additional image processing
 - Beer:
 - For all your friends who look funny on the street taking pictures of traffic signs for the dataset

Traffic Sign Cookbook

- Ingredients:
 - Deep learning library:
 - for defining and training a model and/or additional processing: Knet.jl, Flux.jl
 - Dataset:
 - German Traffic Sign Dataset, cleaned and updated with Romanian traffic signs -> 60 classes and counting
 - WebApp/Server library:
 - Genie.jl
 - Database library:
 - SQLite.jl
 - Server (hardware)
 - Android app:
 - For taking pictures and sending them to the server for storage

Deep Learning Library

- Knet.jl:
 - Implemented in Julia (so is Flux.jl)
 - Straightforward tutorials backed in the package itself
 - Good GPU integration
 - Good memory management

Dataset

- German Traffic Sign Dataset:
 - Advantages:
 - ~40 classes
 - Disadvantages:
 - Some signs are different from the Romanian ones
 - Some important signs are missing
 - Blurred images
 - Dark images
 - Low resolution images
- Fix the shortcomings:
 - Gather more pictures!
 - Apply data augmentation

WebApp/Server Library

- Genie.jl:
 - Implemented in Julia
 - Used for running a BE server which:
 - Receives images from the Phone App
 - Stores images
 - Stores additional data about images in SQLite database
 - Processes images in order to detect traffic signs

Choosing an architecture...

- R-CNN
- Faster R-CNN
- YOLO
- Single Shot Detector



R-CNN

- aka: Region Proposal Conv. Nets
- Two stage processing:
 - Classical image processing:
 - Region Proposal (ROI) by applying image segmentation
 - DL classifier :
 - applied on the ROIs

Processing steps



Segmentation ↓



Crop & process



Apply model ↑

Label & confidence level

```
labels, confidence = segm_img(model, img, 0.15)
("STOP", "0.9978752")
```

Conclusions

- What we achieved (so far):
 - Server: running
 - Dataset: growing
 - Database: in place
 - Image segmentation and region proposal: done
 - DL model: defined from scratch and trained
 - * Android app: functional (used it to send images to server)
- What needs improvement:
 - Better image preprocessing to accommodate different lighting conditions
 - Optimize ROI selection and cropping speed
 - Try other architectures: R-CNN, YOLO (work in progress)



Q&A